

Generative Models

Autoencoder: Encoder



THE UNIVERSITY OF
SOUTHERN MISSISSIPPI®

Generative Models

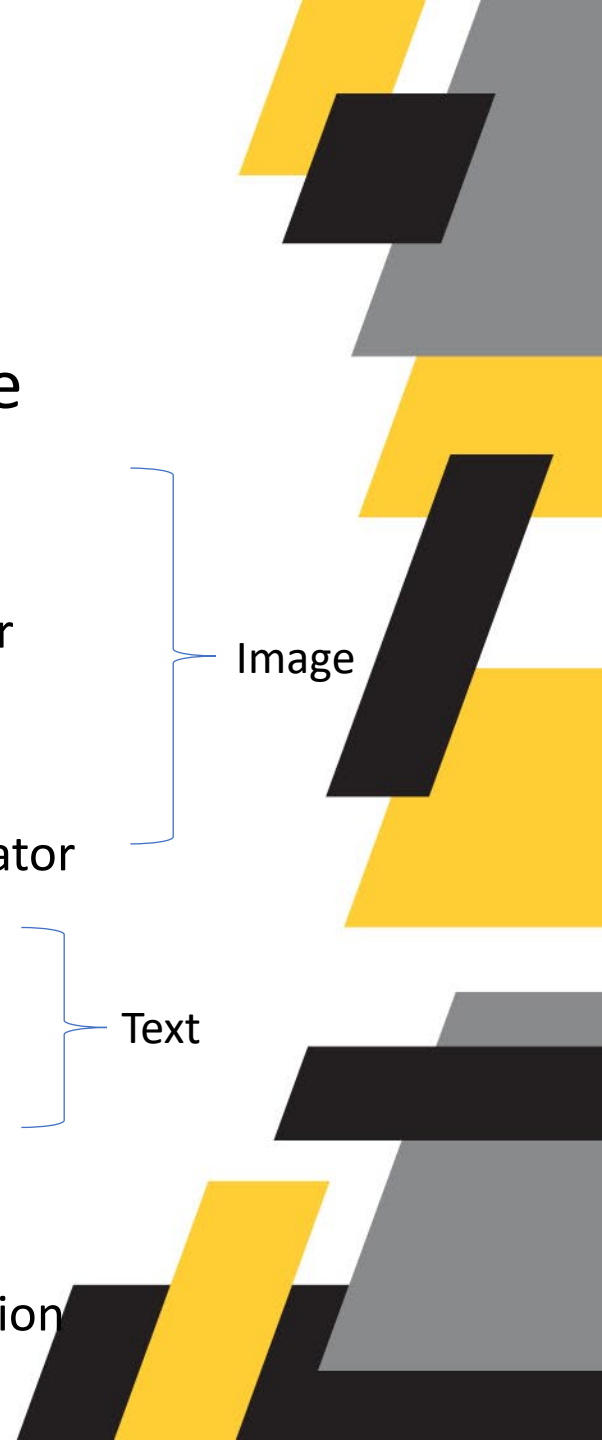
- Is a type of machine learning algorithm for data creation.
 - Data generated is similar to the data used to train a model
- Applications:
 - Image Synthesis
 - Natural language processing
 - Density estimation

• Variations of Generative Models:

- Autoencoders
 - Variational Autoencoder
- Generative Adversarial Networks
 - Generator vs. discriminator
- Autoregressive models
 - LSTM (RNN)
- Transformer
 - GPT
- Diffusion Models
 - Works on noise generation

Image

Text

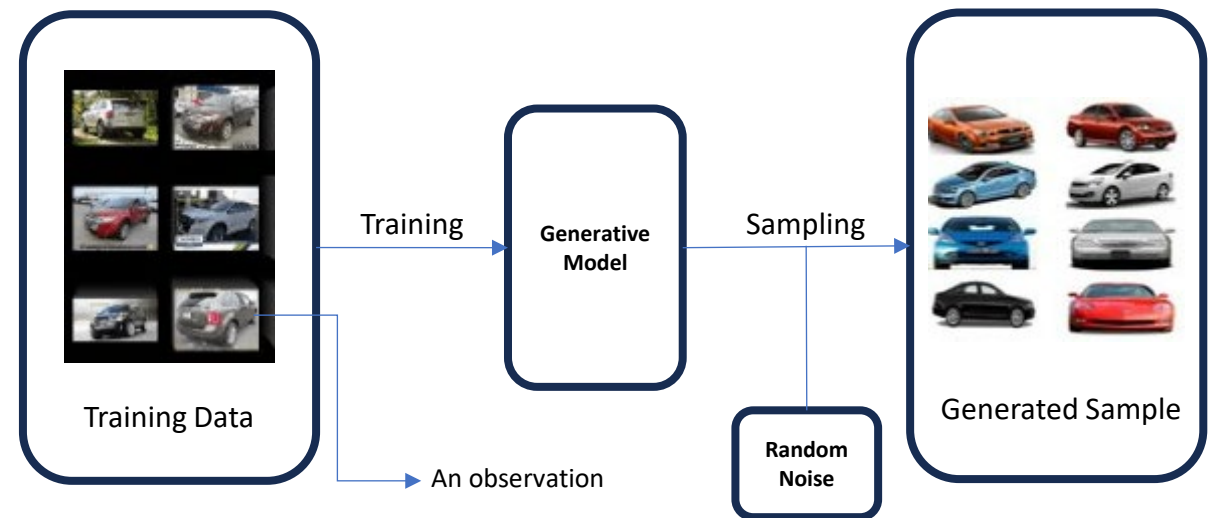


Generative Models (GM)

- Train a generative model to capture rules that govern the relationship between pixels in an image of, for example, horses, cars, human faces, etc.
- Then, sample from this model to create realistic images that did not exist in the original dataset.
- Two main approaches for GM applications:
 - Image
 - Text
- Each observation consists of many features:
 - For image problems: individual pixel values
 - For text problems: individual words or groups of letters

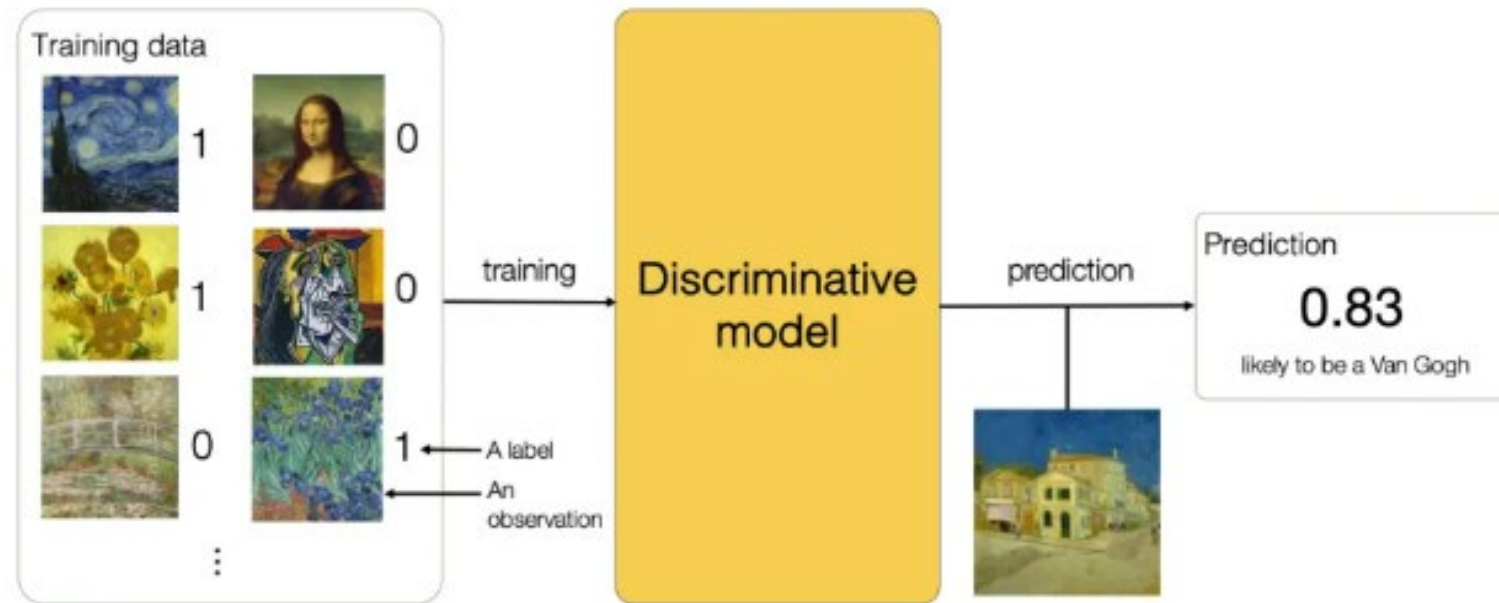
Generative Models (GM)

- GMs are probabilistic rather than deterministic
 - Reason: sample many variations of the output
 - If probabilistic, a random component influences the individual samples generated by the model.
 - Instead of: get the same (or approximated) output every time



Generative vs. Discriminative Modeling

- Discriminative:
 - Each observation in the training data has a *label*
 - For a binary classification, 1 – for Van Gogh, 0 – for non Van Gogh



Generative vs. Discriminative Modeling

- Formally:
 - Discriminative modeling estimates $p(y|x)$
 - Model the probability of a *label* y given some observation x
 - Generative modeling estimates $p(x)$
 - Model the probability of observing an observation x . Sampling from this distribution allows us to generate new observations.
 - Conditional Generative Models: $p(x|y)$
 - The probability of seeing an observation x with a specific label y .
 - For example, the dataset is about different types of vehicles, tell the model to generate an image of an 18-wheeler.

Generative Models (GM): Framework

- We have a dataset of observations X
- We assume that the observations have been generated according to some unknown distribution, P_{data}
- We want to build a GM P_{model} that mimics P_{data} . If we achieve this goal, we can sample from P_{model} to generate observations that appear to have been drawn from P_{data}
- The ideal for P_{model} is:

Generative Models (GM): Framework


- The ideal for P_{model} is:
 - Accuracy:
 - If the accuracy of P_{model} is high for a generated observation, it should look like it has been drawn from P_{data}
 - Generation
 - It should be possible to easily sample a new observation from P_{model}
 - Representation
 - It should be possible to understand how different high-level features in the data are represented by P_{model}

Generative Models (GM): Representation

- Representation of high-dimensional data (representation learning):
 - Instead of trying to model high-dimensional data sample space directly, we describe each observation in the training set using lower-dimensional **latent space** and map it to a point in the original domain.
 - “Short definition”: each point in the **latent space** is a representation of some high-dimensional observation

Generative Models (GM): Autoencoder

- Autoencoder:
 - Short definition: A neural network that is trained to perform the task of encoding and decoding an item
 - The output should be as close to the original item as possible




Generative Models (GM): Autoencoder

- Why reconstruct images that are already available?
 - Sampling from the latent space -> creation (generation) of new unseen image
- Embedding:
 - Compression of the original image into a lower-dimensional latent space
 - Done by the **Encoder**
- **Encoder**
 - Main job: to take an input image and map it to an embedding vector in the latent space

Generative Models (GM): Autoencoder: Encoder

- Encoder architecture:
 - Define the Input layer of the encoder (the image)
 - Stack Conv2D layers sequentially on top of each other
 - Remember the convolutional process for images:
 - Some pixel data will be lost due to filtering!
 - Flatten the last convolutional layer to a vector
 - **Connect this vector to the 2D embedding with a Dense Layer**
 - This allows to plot the latent space



For example: a sample can be: $[-2.5, -6.0]$

● Generated image mapped to latent space

Generative Models


Autoencoder: Decoder



THE UNIVERSITY OF
SOUTHERN MISSISSIPPI®

Generative Models (GM): Autoencoder: Decoder


- The decoder is a mirror image of the encoder:
 - Instead of convolutional layers, the decoder uses **convolutional transpose layers**.
 - Convolutional transpose layers:
 - In **regular convolutional layers**, the resulting output is reduced due to the filtering processes, remember key terms:
 - For example, an input image of $32 \times 32 \times 1$ with stride = 2 will produce a $16 \times 16 \times 1$ output
 - *Stride: step size used by the layer to move the filters across the input. By Increasing the stride, the size of the output (tensor) is reduced.
 - *Padding: “pads” the input data with zeros (zero padding) so that the output size from the layer is exactly the same as the input size when strides = 1



A $3 \times 3 \times 1$ kernel (gray) being passed over a $5 \times 5 \times 1$ input image (blue), with padding = 'same' and strides = 1 to generate the $5 \times 5 \times 1$ output (green).
Padding = 'same' allows to easily keep track of the size of the tensor.

Generative Models (GM): Autoencoder: Decoder

- In the case of **convolutional transpose layers**:
 - Use the same principle as a standard convolutional layer, but setting **strides = 2** doubles the size of the input tensor in both dimensions.
 - In convolutional transpose layers, the strides parameter determines the internal zero padding between pixels in the image



A 3x3x1 filter (gray) is being passed across a 3x3x1 image (blue) with strides = 2 to produce a 6x6x1 output sensor (green).




The **Conv2DTranspose** method in Keras allows the conv. transpose operations until the image reaches its original dimensions. The 0-padding method in transpose “adds” extra space according to the stride parameter.

Generative Models (GM): Autoencoder: Joining the Encoder and Decoder

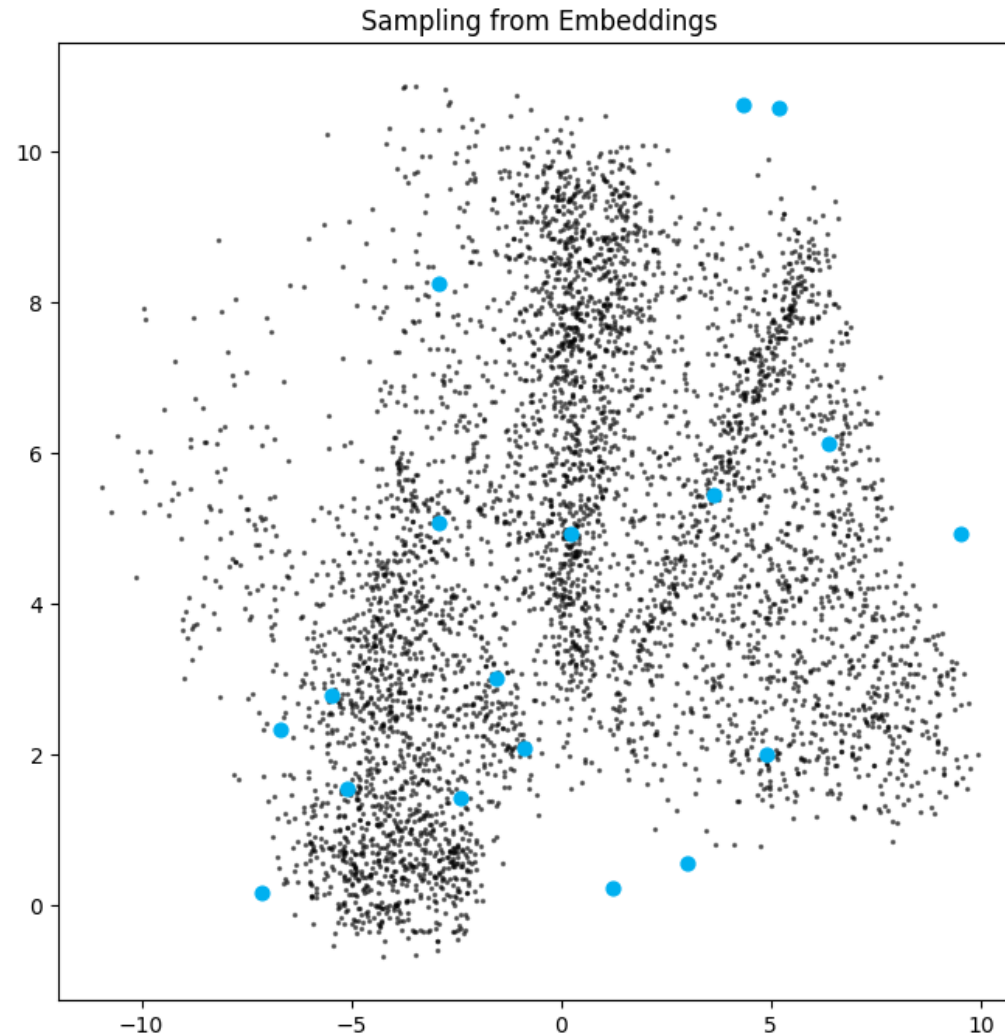
- To train the encoder and decoder simultaneously we need to define a model that will represent the flow of image
 - Flow: Encoder and back out through the decoder
- In Keras, the flow can be specified:
 - Output from the autoencoder is:
 - The output from the encoder after it has been passed through the decoder.
- `autoencoder = Model(encoder_input, decoder(encoder_output))`
- Model takes an image and passes it through the encoder and back out through the decoder to generate a reconstruction of the original image.

Generative Models (GM): Decoder

- Generating new images
 - We can generate novel images by sampling some points in the latent space and using the decoder to convert these back into pixel space.



Generative Models (GM): Decoder



Example

On Canvas: [Download Autoencoder_test.py](#)



THE UNIVERSITY OF
SOUTHERN MISSISSIPPI.