


Introduction to Neural Networks



THE UNIVERSITY OF
SOUTHERN MISSISSIPPI®

In RL

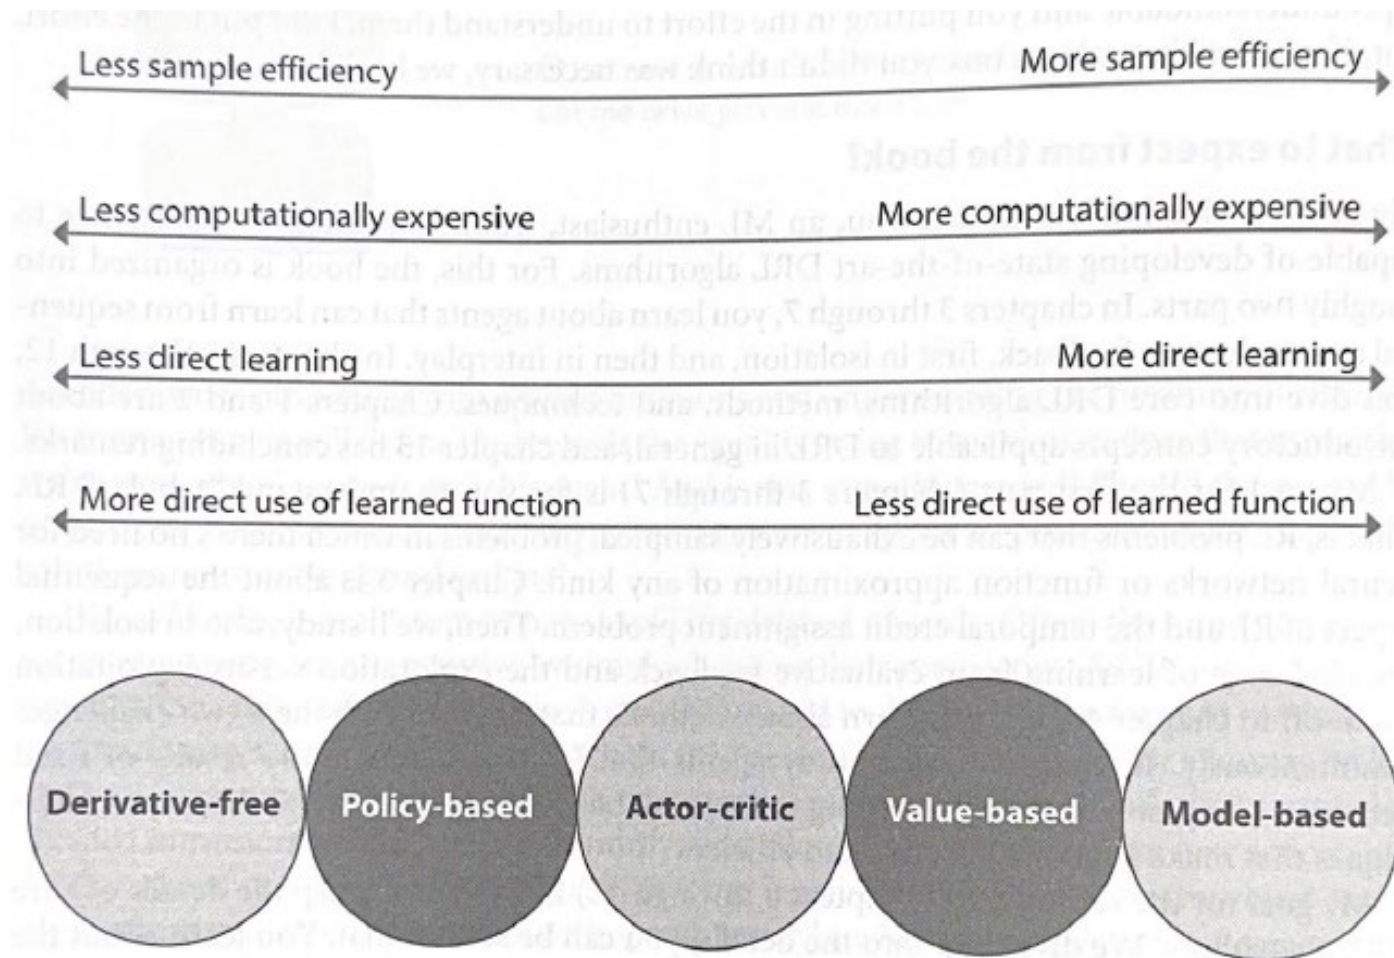
- Three approaches to learning:
 - Policies: map observations to actions
 - Depending on configuration:
 - Off-policy: Q-Learning
 - On-policy: SARSA
 - Models: learn the environment based on a model
 - Target model: goal of the agent
 - Base model: interaction of the agent with the environment (computationally intensive)
 - Value functions: learn to estimate the reward-to-go on mappings



- As more observations, more Q-values, more new states are generated:
 - Interaction of the RL agent and environment grows significantly.
 - An RL approach for more complex environments will be computationally expensive.
- Since RL is based on “trial – error”, an RL agent will have difficulties in more complex environments
 - The “prediction” section of Q-learning
 - Also based on “trial – error” by applying (state, action) from previous observations
- To solve this, implement ***approximations***:
 - Neurons




- More resources are needed depending on approach:



Neuron

- A neuron consists of:
 - An input, **u**
 - An output, **y**
 - An activation signal
 - A threshold, **θ**
 - A weight, **w**
 - *bias, **b**, can be part of the neuron, it can also be considered a hyperparameter





- Mathematically:

$$y = f(x), x = \sum_{i=1}^n w_i u_i$$

Where: x is a weighted sum, and $f(x)$ an activation function depending on the parameter θ

Neuron

- When several neurons are configured in layers, and all neurons are connected to each other, it's known as a:
 - Fully connected neural network
 - In the simplest form, this is called: **Artificial Neural Network, ANN**
 - A model using an ANN with multiple (usually more than 2 or 3) NN layers is referred to as a model using *Deep Learning*.



Neural network

Layers

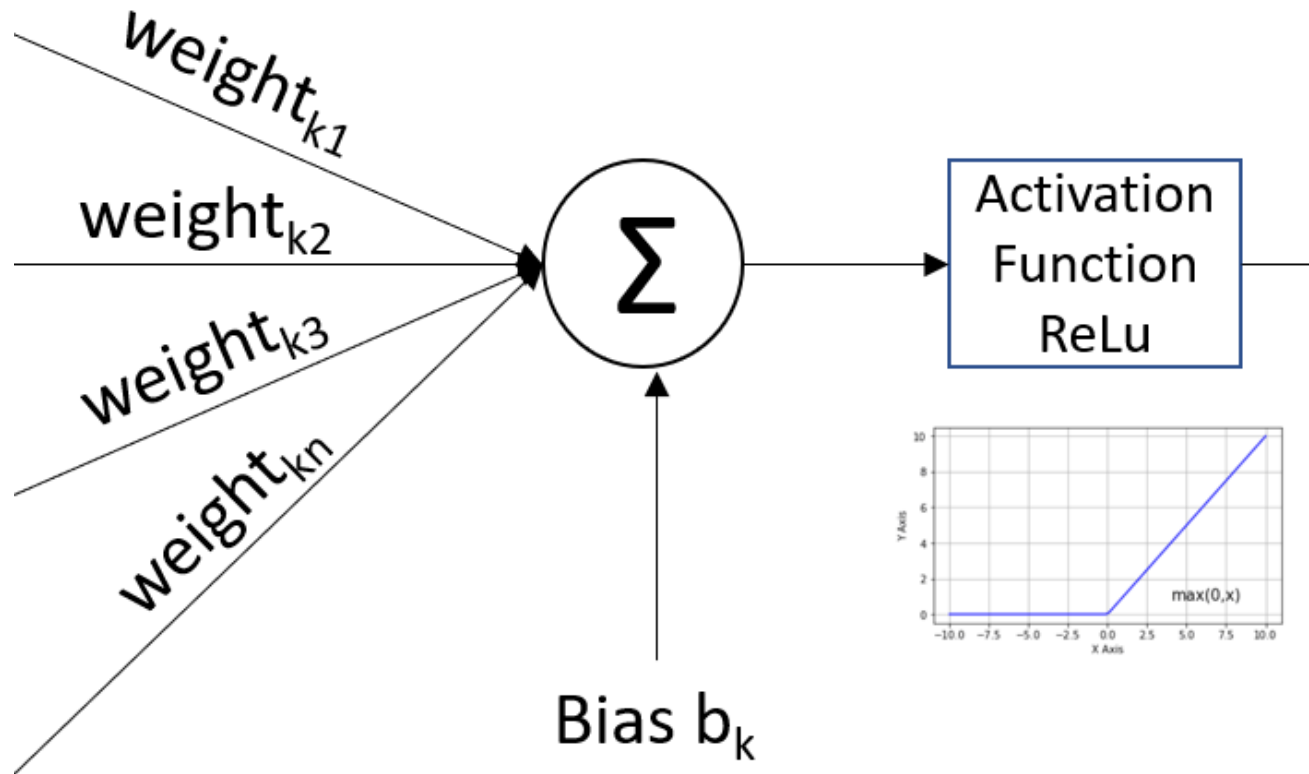
- In an ANN, in its basic structure: three layers
 - Input Layer:
 - Data coming from “the environment”
 - In DRL these will be observations from the agent
 - In Deep Learning:
 - Labeled data, i.e.
 - Depending on application:
 - Convolutional (& pooling layer): format the input data
 - Convolutional Neural Network, CNN
 - Recurrent: memory-based implementation
 - Recurrent Neural Network, RNN
 - Hidden Layer(s)
 - Where all approximations occur
 - Output Layer
 - Results of the approximations



Approximation

- In one neuron, calculations are occurring:
- Neurons in the hidden layer:
 - Input: weighted sum of neurons in previous layers
 - Weighted sum: $\text{input} * \text{weight} + \text{bias}$ (bias can be optional, however, its is most likely to be included in the calculation)
 - Activation: check if the final weighted sum is within threshold parameters
 - If final weighted sum is within parameters:
 - Neuron remains active
 - If not:
 - Send weighted sum to next layer
 - Deactivate neuron
 - Neuron is no longer used on the next episode (epoch)
- Output layer:
 - Results of approximations:
 - transformed into formatted data

Hidden Layer





Activation

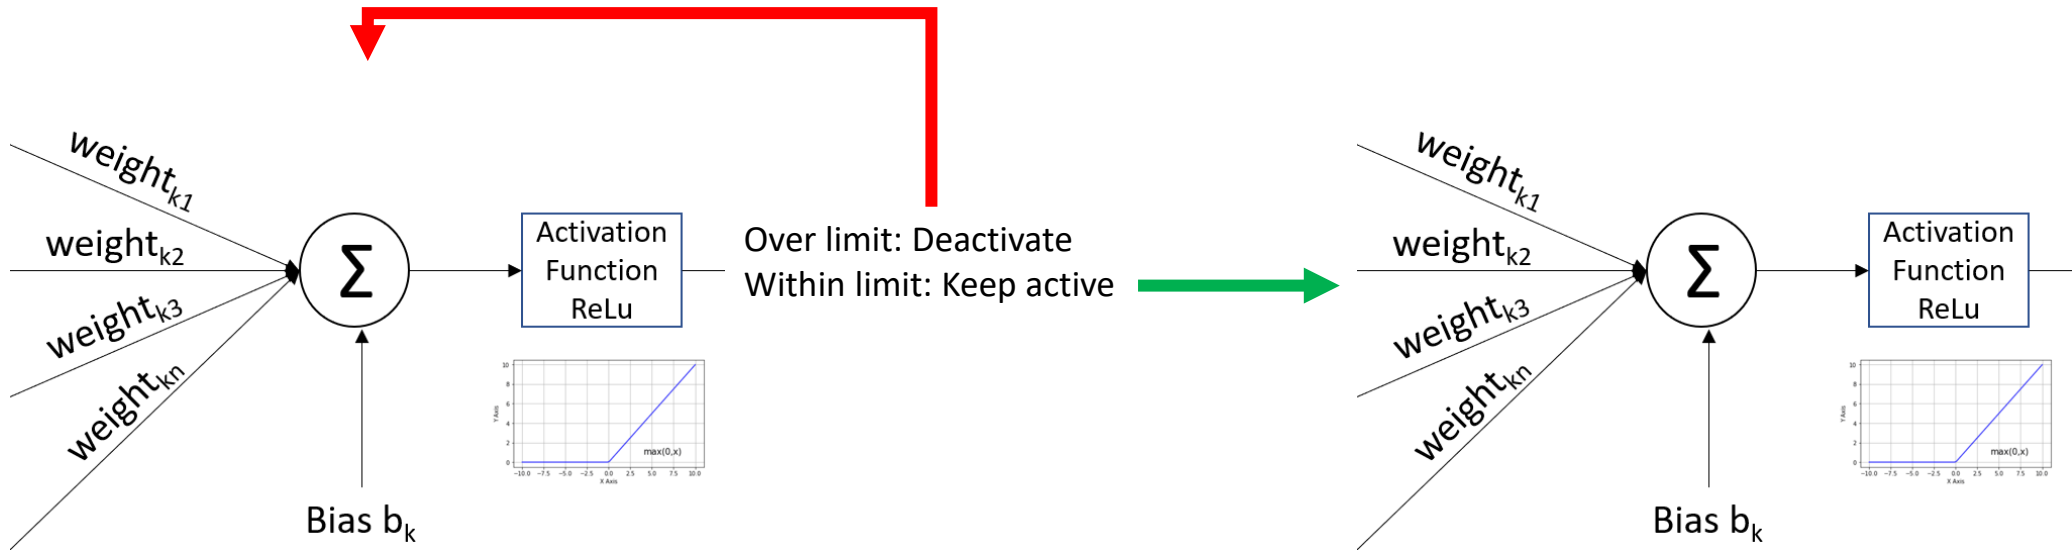
- Activation is used to decide whether the neuron in question should be kept active or not:

- Methods:

- Softmax: based on probabilities
 - Calculates a probability
- Sigmoid: when the output needs to
 - based on probabilities
- Tanh: similar to sigmoid, except that
 - Tanh will accept negatives
- Rectified Linear Unit (ReLU):
 - $f(x)$ is 0 when x is 0



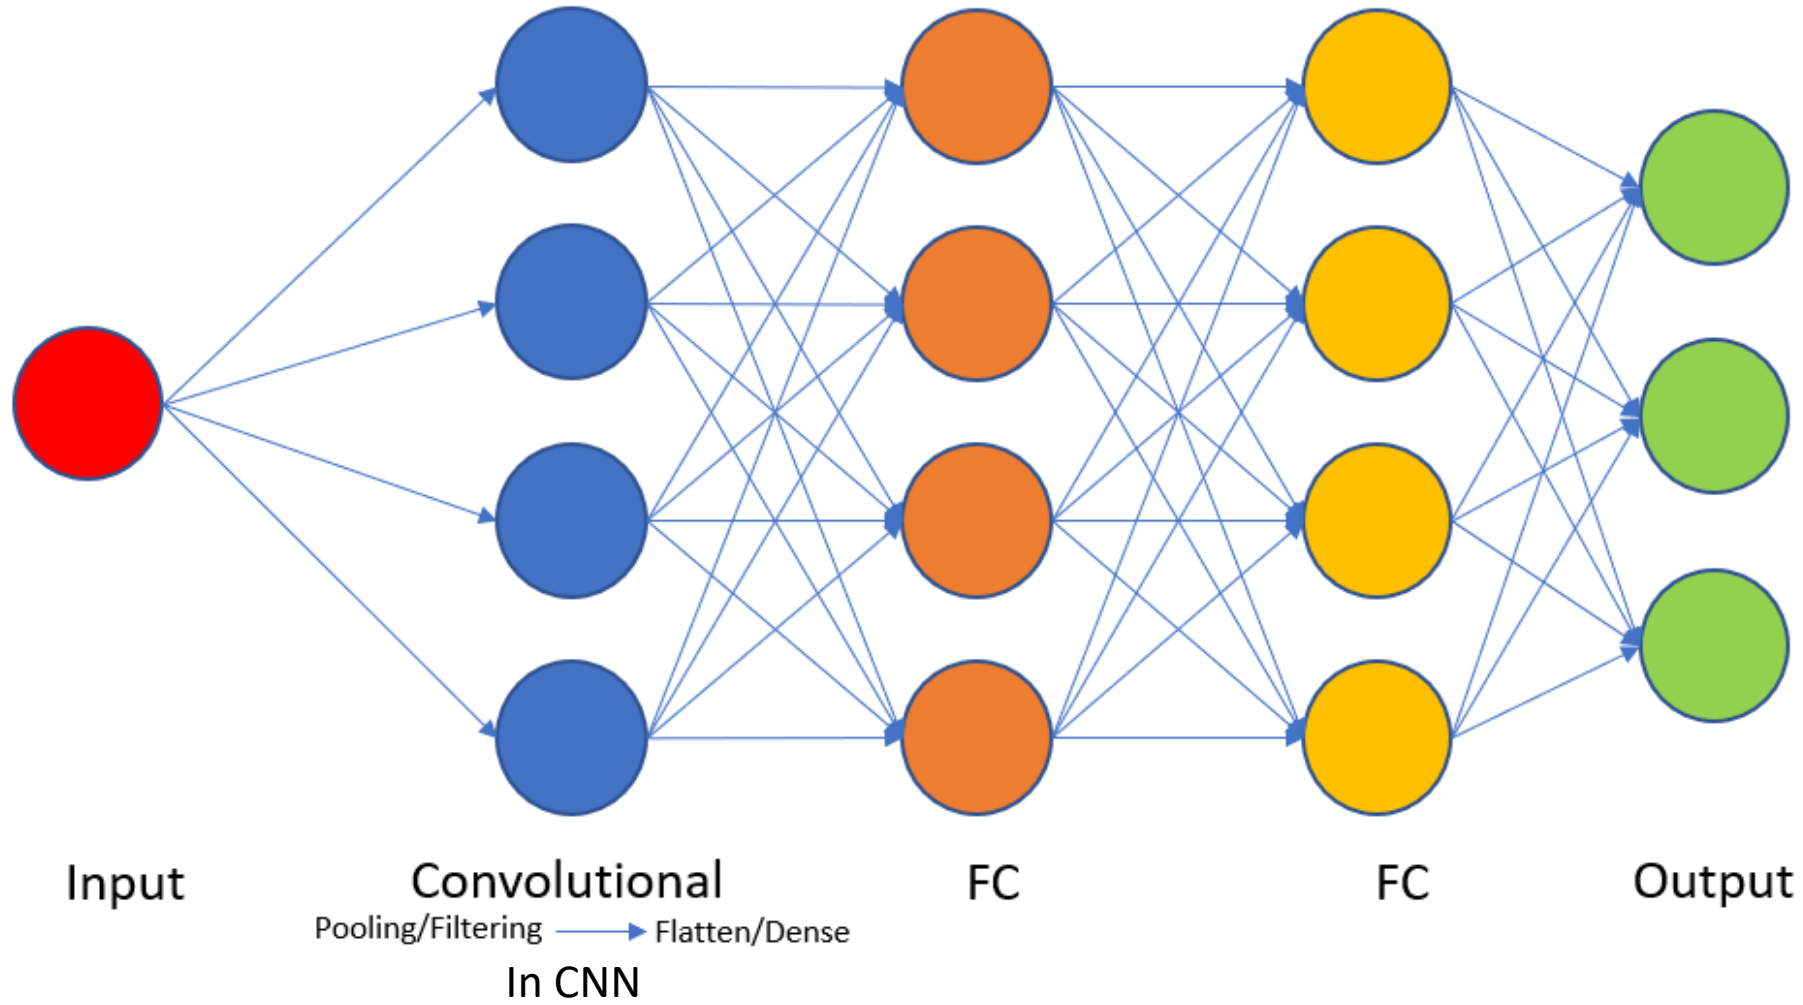
Activation



- Before approximations begin:
 - Data needs to be “transformed” into proper neural input
 - Convolutional layer (& pooling layer):
 - Filtering
 - Pooling
 - Flatten & dense



Representation of a fully connected ANN





THE UNIVERSITY OF
SOUTHERN MISSISSIPPI.